

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error

# Load the text file
df = pd.read_csv('household_power_consumption.txt', sep=';', parse_dates={'datetime': ['Date', 'Time']}, infer_datetime_format=True, low_memory=True)

# Fill missing values with the mean
df['Global_active_power'] = pd.to_numeric(df['Global_active_power'], errors='coerce')
df['Global_active_power'].fillna(df['Global_active_power'].mean(), inplace=True)

# Set the datetime as the index
df.set_index('datetime', inplace=True)

# Resample the data to daily frequency
df_daily = df.resample('D').sum()

# Basic statistics
print(df_daily.describe())

# Check for missing values
print(df_daily.isnull().sum())


# Plotting the data
plt.figure(figsize=(14, 7))
plt.plot(df_daily.index, df_daily['Global_active_power'])
plt.title('Global Active Power over Time')
plt.xlabel('Date')
plt.ylabel('Global Active Power (kilowatts)')
plt.show()
```

```

<ipython-input-15-3570c6b41cfa>:9: FutureWarning: The argument 'infer_datetime_format'
df = pd.read_csv('household_power_consumption.txt', sep=';', parse_dates={'datetime':
<ipython-input-15-3570c6b41cfa>:9: UserWarning: Parsing dates in %d/%m/%Y %H:%M:%S form
df = pd.read_csv('household_power_consumption.txt', sep=';', parse_dates={'datetime':
      Global_active_power  Global_reactive_power  Voltage \
count      1442.000000      1442.000000      1442.000000
mean       1571.001338      175.815258      342266.507732
std        595.405647      51.998109      36707.752471
min         250.298000       0.000000       0.000000
25%        1191.182615      140.911500      345621.842500
50%        1559.085000      170.290000      346934.710000
75%        1889.859500      202.372500      348251.527500
max         4773.386000      417.834000      356306.410000

      Global_intensity  Sub_metering_1  Sub_metering_2  Sub_metering_3
count      1442.000000      1442.000000      1442.000000      1442.000000
mean       6576.681415      1594.407074      1845.375173      9178.340499
std       2559.505974      1587.840580      2089.590342      3787.898093
min         0.000000       0.000000       0.000000       0.000000
25%        4988.800000      555.500000      424.250000      6604.250000
50%        6510.300000      1109.000000      678.500000      9251.000000
75%        7953.350000      2196.750000      2712.750000      11708.500000
max       20200.400000     11178.000000     12109.000000     23743.000000
Global_active_power      0
Global_reactive_power    0
Voltage                  0
Global_intensity        0
Sub_metering_1          0
Sub_metering_2          0
Sub_metering_3          0
dtype: int64

```



```

# Define the training and testing periods
# Adjust these dates to have some overlap and avoid a gap
train_end_date = '2009-12-31'
test_start_date = '2009-01-01'

# Split the data
train = df_daily[:train_end_date]
test = df_daily[test_start_date:]

# Check if the test set is empty and handle the case
if test.empty:
    print("Warning: Test set is empty. Adjust the date range.")
else:
    # Fit an ARIMA model
    # You might need to experiment with different (p, d, q) orders

```

```


model = ARIMA(train['Global_active_power'], order=(5, 1, 0))
model_fit = model.fit()

# Forecasting
forecast = model_fit.forecast(steps=len(test))
test['Forecast'] = forecast.values

# Calculate MSE
mse = mean_squared_error(test['Global_active_power'], test['Forecast'])
print(f'Mean Squared Error: {mse}')

# Plot the results
plt.figure(figsize=(14, 7))
plt.plot(train.index, train['Global_active_power'], label='Training Data')
plt.plot(test.index, test['Global_active_power'], label='Test Data')
plt.plot(test.index, test['Forecast'], label='Forecast')
plt.legend(loc='upper left')
plt.title('Global Active Power Forecast')
plt.xlabel('Date')
plt.ylabel('Global Active Power (kilowatts)')
plt.show()

```

 <ipython-input-19-40b94adc05a4>:21: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/user>

```

test['Forecast'] = forecast.values
Mean Squared Error: 689430.7539605086

```

